

Hands_on_Handling_missing_value_with_List_&_Pairwise_Deletion_with

September 1, 2023

0.1 Handling Missing Value in Machine Learning

0.1.1 Finding the missing values in Dataset

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import matplotlib inline
```

```
[2]: # Import the dataset
df_titanic = pd.read_csv("https://raw.githubusercontent.com/atulpatelDS/
↳Data_Files/master/Titanic/titanic_train.csv")
```

```
[3]: df_titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     891 non-null   int64
 1   Survived        891 non-null   int64
 2   Pclass         891 non-null   int64
 3   Name           891 non-null   object
 4   Sex            891 non-null   object
 5   Age            714 non-null   float64
 6   SibSp          891 non-null   int64
 7   Parch         891 non-null   int64
 8   Ticket         891 non-null   object
 9   Fare          891 non-null   float64
10   Cabin         204 non-null   object
11   Embarked      889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[4]: df_titanic.describe(include=['O']) ## For categorical
```

```
[4]:
```


	Name	Sex	Ticket	Cabin	Embarked
count	891	891	891	204	889
unique	891	2	681	147	3
top	Braund, Mr. Owen Harris	male	347082	B96 B98	S
freq	1	577	7	4	644

```
[5]: df_titanic.isnull().sum()
```

```
[5]: PassengerId      0
      Survived        0
      Pclass          0
      Name            0
      Sex             0
      Age            177
      SibSp           0
      Parch           0
      Ticket          0
      Fare            0
      Cabin           687
      Embarked        2
      dtype: int64
```


```
[7]: # Program to visualize missing values in dataset
      # Importing the libraries
      import missingno as msno
      # Visualize missing values as a matrix
      msno.matrix(df_titanic,figsize=(12,6))
```

```
[7]: <Axes: >
```




```
[8]: # Visualize the number of missing
# values as a bar chart
msno.bar(df_titanic,figsize=(10,4))
```

[8]: <Axes: >



```
[9]: #sns.heatmap(df_titanic.isnull(),cmap="viridis")
plt.figure(figsize=(10,6))
sns.heatmap(
    data=df_titanic.isnull(),
    cmap=sns.color_palette(['black', 'yellow', 'orange', 'white']),cbar=False)
```

[9]: <Axes: >



It is clearly visible that most of the null values are available in column AGE and CABIN

```
[10]: # lets find out the percentage of missing value in each column
Percent_Missing_Value = df_titanic.isnull().sum()*100/len(df_titanic)
Percent_Missing_Value
```

```
[10]: PassengerId    0.000000
Survived           0.000000
Pclass             0.000000
Name               0.000000
Sex                0.000000
Age                19.865320
SibSp              0.000000
Parch              0.000000
Ticket             0.000000
```

```
Fare          0.000000
Cabin         77.104377
Embarked      0.224467
dtype: float64
```

0.2 Method 1- Removal or Deletion of missing value

List wise Deletion of Missing Value

- Use dropna(), drop() functions

```
[11]: # Import the dataset
df_titanic = pd.read_csv("https://raw.githubusercontent.com/atulpatelDS/
↳Data_Files/master/Titanic/titanic_train.csv")
```

```
[12]: df_titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[13]: df_titanic.isnull().sum()
```

```
[13]: PassengerId     0
Survived            0
Pclass              0
Name                0
Sex                 0
Age                177
SibSp              0
Parch              0
Ticket             0
Fare               0
```

```
Cabin          687
Embarked       2
dtype: int64
```

```
[14]: ## Total no of samples where any column has NaN
df_titanic[(df_titanic["Age"].isnull()==True)|
            (df_titanic["Cabin"].isnull()==True)|
            (df_titanic["Embarked"].isnull()==True)]
```

```
[14]:
```

	PassengerId	Survived	Pclass	Name \
0	1	0	3	Braund, Mr. Owen Harris
2	3	1	3	Heikkinen, Miss. Laina
4	5	0	3	Allen, Mr. William Henry
5	6	0	3	Moran, Mr. James
7	8	0	3	Palsson, Master. Gosta Leonard
..
884	885	0	3	Sutehall, Mr. Henry Jr
885	886	0	3	Rice, Mrs. William (Margaret Norton)
886	887	0	2	Montvila, Rev. Juozas
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
890	891	0	3	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	male	22.0	1	0	A/5 21171	7.2500	NaN	S
2	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
4	male	35.0	0	0	373450	8.0500	NaN	S
5	male	NaN	0	0	330877	8.4583	NaN	Q
7	male	2.0	3	1	349909	21.0750	NaN	S
..
884	male	25.0	0	0	SOTON/OQ 392076	7.0500	NaN	S
885	female	39.0	0	5	382652	29.1250	NaN	Q
886	male	27.0	0	0	211536	13.0000	NaN	S
888	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
890	male	32.0	0	0	370376	7.7500	NaN	Q

[708 rows x 12 columns]

```
[15]: df_titanic.shape
```

```
[15]: (891, 12)
```

```
[16]: ## Lets delete all the Rows where we have NaN values
df_titanic_new = df_titanic.dropna()
```

```
[17]: df_titanic_new.shape
```

```
[17]: (183, 12)
```

```
[18]: 891-708
```

```
[18]: 183
```

We have total 891 samples and 708 out of 891 have the NaN value. If we delete all the rows where we have NaN value then we can only get sample size 183, which is very less value to make the model so we cannot delete all the rows.

```
[19]: df_titanic_new[(df_titanic_new["Age"].isnull()==True)|
                    (df_titanic_new["Cabin"].isnull()==True)|
                    (df_titanic_new["Embarked"].isnull()==True)]
```

```
[19]: Empty DataFrame
      Columns: [PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket,
      Fare, Cabin, Embarked]
      Index: []
```

```
[20]: df_titanic_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 183 entries, 1 to 889
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      183 non-null   int64
1   Survived         183 non-null   int64
2   Pclass           183 non-null   int64
3   Name             183 non-null   object
4   Sex              183 non-null   object
5   Age              183 non-null   float64
6   SibSp            183 non-null   int64
7   Parch            183 non-null   int64
8   Ticket           183 non-null   object
9   Fare             183 non-null   float64
10  Cabin            183 non-null   object
11  Embarked         183 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 18.6+ KB
```

```
[21]: df_titanic[(df_titanic["Age"].isnull()==True)&
               (df_titanic["Cabin"].isnull()==True)
               ]
```

```
[21]:
```

	PassengerId	Survived	Pclass	Name \
5	6	0	3	Moran, Mr. James
17	18	1	2	Williams, Mr. Charles Eugene
19	20	1	3	Masselmani, Mrs. Fatima
26	27	0	3	Emir, Mr. Farred Chehab

28	29	1	3	O'Dwyer, Miss. Ellen "Nellie"
..
859	860	0	3	Razi, Mr. Raihed
863	864	0	3	Sage, Miss. Dorothy Edith "Dolly"
868	869	0	3	van Melkebeke, Mr. Philemon
878	879	0	3	Laleff, Mr. Kristo
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
5	male	NaN	0	0	330877	8.4583	NaN	Q
17	male	NaN	0	0	244373	13.0000	NaN	S
19	female	NaN	0	0	2649	7.2250	NaN	C
26	male	NaN	0	0	2631	7.2250	NaN	C
28	female	NaN	0	0	330959	7.8792	NaN	Q
..
859	male	NaN	0	0	2629	7.2292	NaN	C
863	female	NaN	8	2	CA. 2343	69.5500	NaN	S
868	male	NaN	0	0	345777	9.5000	NaN	S
878	male	NaN	0	0	349217	7.8958	NaN	S
888	female	NaN	1	2	W./C. 6607	23.4500	NaN	S

[158 rows x 12 columns]

We can see there are 158 Rows where we have NaN value in both Age and Cabin in same row. We can delete these rows.

```
[22]: # lets drop those rows where both Age and Cabin are null
df_titanic_new1 = df_titanic.drop(df_titanic.index[(df_titanic["Age"].
↪isnull()==True)&
(df_titanic["Cabin"].isnull()==True)
])
```

```
[23]: df_titanic_new1
```

```
[23]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
..	
885	886	0	3	
886	887	0	2	
887	888	1	1	
889	890	1	1	
890	891	0	3	

Name	Sex	Age	SibSp	\
------	-----	-----	-------	---

0		Braund, Mr. Owen Harris	male	22.0	1
1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1
2		Heikkinen, Miss. Laina	female	26.0	0
3		Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1
4		Allen, Mr. William Henry	male	35.0	0
..	
885		Rice, Mrs. William (Margaret Norton)	female	39.0	0
886		Montvila, Rev. Juozas	male	27.0	0
887		Graham, Miss. Margaret Edith	female	19.0	0
889		Behr, Mr. Karl Howell	male	26.0	0
890		Dooley, Mr. Patrick	male	32.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..
885	5	382652	29.1250	NaN	Q
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[733 rows x 12 columns]

```
[24]: 891-158
```

```
[24]: 733
```

```
[25]: # If we want to drop all those rows where we have all NaN value in all columns
      ↪ in a row.
df_titanic_new3 = df_titanic.dropna(how="all")
```

```
[26]: df_titanic_new3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age         714 non-null    float64
```

```

6   SibSp      891 non-null   int64
7   Parch      891 non-null   int64
8   Ticket     891 non-null   object
9   Fare       891 non-null   float64
10  Cabin      204 non-null   object
11  Embarked    889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

As we don't have any row where all the value is NaN. So It did not delete any row.

```

[27]: # Import Dataset
import pandas as pd
df_sales = pd.read_excel("https://raw.githubusercontent.com/atulpatelDS/
↳Data_Files/master/Feature_Engineering/Missing_Value/Sales.xlsx")

```

```

[28]: df_sales

```

```

[28]:
      Date  Day_Temp  No_of_Customers  Sales
0  2020-10-01      30.0             100.0  3112.0
1  2020-10-02       NaN             115.0  3682.0
2  2020-10-03      31.0              NaN  2774.0
3  2020-10-04      29.0             105.0  3182.0
4  2020-10-05      33.0             104.0  1368.0
5  2020-10-07       NaN              NaN     NaN
6  2020-11-24      26.0             90.0  4232.0
7  2020-11-25       NaN             96.0     NaN
8  2020-11-26      27.0            100.0  2356.0
9  2020-11-28       NaN              NaN     NaN
10 2020-11-29      23.0             94.0  1254.0
11 2020-11-30      22.0             91.0  4232.0

```

```

[29]: df_sales.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12 entries, 0 to 11
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  12 non-null    datetime64[ns]
1   Day_Temp              8 non-null     float64
2   No_of_Customers       9 non-null     float64
3   Sales                 9 non-null     float64
dtypes: datetime64[ns](1), float64(3)
memory usage: 512.0 bytes

```

```

[30]: ## Set Date as Index
df_sales.set_index("Date",inplace=True)

```

```
[31]: df_sales
```

```
[31]:
```

	Day_Temp	No_of_Customers	Sales
Date			
2020-10-01	30.0	100.0	3112.0
2020-10-02	NaN	115.0	3682.0
2020-10-03	31.0	NaN	2774.0
2020-10-04	29.0	105.0	3182.0
2020-10-05	33.0	104.0	1368.0
2020-10-07	NaN	NaN	NaN
2020-11-24	26.0	90.0	4232.0
2020-11-25	NaN	96.0	NaN
2020-11-26	27.0	100.0	2356.0
2020-11-28	NaN	NaN	NaN
2020-11-29	23.0	94.0	1254.0
2020-11-30	22.0	91.0	4232.0

```
[32]: ## Lets Delete row where atleast one NaN value
df_sales_1 = df_sales.dropna()
```

```
[33]: df_sales_1
```

```
[33]:
```

	Day_Temp	No_of_Customers	Sales
Date			
2020-10-01	30.0	100.0	3112.0
2020-10-04	29.0	105.0	3182.0
2020-10-05	33.0	104.0	1368.0
2020-11-24	26.0	90.0	4232.0
2020-11-26	27.0	100.0	2356.0
2020-11-29	23.0	94.0	1254.0
2020-11-30	22.0	91.0	4232.0

```
[34]: ## If we want to delete Row where all rows data are NaN.
df_sales_2 = df_sales.dropna(how="all")
```

```
[35]: df_sales_2
```

```
[35]:
```

	Day_Temp	No_of_Customers	Sales
Date			
2020-10-01	30.0	100.0	3112.0
2020-10-02	NaN	115.0	3682.0
2020-10-03	31.0	NaN	2774.0
2020-10-04	29.0	105.0	3182.0
2020-10-05	33.0	104.0	1368.0
2020-11-24	26.0	90.0	4232.0
2020-11-25	NaN	96.0	NaN
2020-11-26	27.0	100.0	2356.0
2020-11-29	23.0	94.0	1254.0

2020-11-30	22.0	91.0	4232.0
------------	------	------	--------

```
[36]: # If we want to keep only those rows where we have atleast one valid value.
# thresh : int, optional : Require that many non-NA values.
df_sales_3 = df_sales.dropna(thresh=1)
df_sales_3
```

```
[36]:
```

	Day_Temp	No_of_Customers	Sales
Date			
2020-10-01	30.0	100.0	3112.0
2020-10-02	NaN	115.0	3682.0
2020-10-03	31.0	NaN	2774.0
2020-10-04	29.0	105.0	3182.0
2020-10-05	33.0	104.0	1368.0
2020-11-24	26.0	90.0	4232.0
2020-11-25	NaN	96.0	NaN
2020-11-26	27.0	100.0	2356.0
2020-11-29	23.0	94.0	1254.0
2020-11-30	22.0	91.0	4232.0

```
[37]: # If we want to keep only those rows where we have atleast two valid value.
# thresh : int, optional : Require that many non-NA values.
df_sales_4 = df_sales.dropna(thresh=2)
df_sales_4
```

```
[37]:
```

	Day_Temp	No_of_Customers	Sales
Date			
2020-10-01	30.0	100.0	3112.0
2020-10-02	NaN	115.0	3682.0
2020-10-03	31.0	NaN	2774.0
2020-10-04	29.0	105.0	3182.0
2020-10-05	33.0	104.0	1368.0
2020-11-24	26.0	90.0	4232.0
2020-11-26	27.0	100.0	2356.0
2020-11-29	23.0	94.0	1254.0
2020-11-30	22.0	91.0	4232.0

```
[38]: # If we want to keep only those rows where we have atleast three valid value.
# thresh : int, optional : Require that many non-NA values.
df_sales_5 = df_sales.dropna(thresh=3)
df_sales_5
```

```
[38]:
```

	Day_Temp	No_of_Customers	Sales
Date			
2020-10-01	30.0	100.0	3112.0
2020-10-04	29.0	105.0	3182.0
2020-10-05	33.0	104.0	1368.0

2020-11-24	26.0	90.0	4232.0
2020-11-26	27.0	100.0	2356.0
2020-11-29	23.0	94.0	1254.0
2020-11-30	22.0	91.0	4232.0

Pair wise Deletion of Missing Value

```
[39]: # Lets Delete the NaN value column wise.
# lets find out the percentage of missing vale in each column
Percent_Missing_Value = df_titanic.isnull().sum()*100/len(df_titanic)
Percent_Missing_Value
```

```
[39]: PassengerId    0.000000
Survived          0.000000
Pclass           0.000000
Name             0.000000
Sex              0.000000
Age             19.865320
SibSp            0.000000
Parch            0.000000
Ticket           0.000000
Fare             0.000000
Cabin            77.104377
Embarked         0.224467
dtype: float64
```

```
[40]: df_titanic.head()
```

```
[40]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1
4             5         0         3

                                     Name    Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                        Heikkinen, Miss. Laina  female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4                        Allen, Mr. William Henry    male  35.0      0

   Parch    Ticket   Fare Cabin Embarked
0      0   A/5 21171   7.2500   NaN        S
1      0   PC 17599  71.2833   C85        C
2      0  STON/O2. 3101282   7.9250   NaN        S
3      0    113803  53.1000  C123        S
4      0    373450   8.0500   NaN        S
```

We can clearly see that only 20% of the AGE data is missing. The proportion of the AGE missing is likely small enough for reasonable replacement with some form of imputation. Now see the CABIN column, It looks like we are missing too much of that data to do something useful with at a basic level. We will drop it.

If columns have more than half of rows as null then the entire column can be dropped.

```
[41]: # Lets delete the Cabin column
df_titanic_new2 = df_titanic.drop("Cabin",axis=1)
```

```
[42]: df_titanic_new2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Embarked        889 non-null   object
dtypes: float64(2), int64(5), object(4)
memory usage: 76.7+ KB
```

```
[43]: # Import Dataset
import pandas as pd
df_saless = pd.read_excel("https://raw.githubusercontent.com/atulpatelDS/
↳Data_Files/master/Feature_Engineering/Missing_Value/Saless.xlsx")
```

```
[44]: df_saless
```

```
[44]:
```

	Date	Store_Type	City_Type	Day_Temp	No_of_Customers	Sales	\
0	2020-10-01	1	1	30.0	100.0	3112.0	
1	2020-10-02	2	1	32.0	115.0	3682.0	
2	2020-10-03	3	3	31.0	NaN	2774.0	
3	2020-10-04	1	2	29.0	105.0	3182.0	
4	2020-10-05	1	2	33.0	104.0	1368.0	
5	2020-10-07	2	2	NaN	NaN	NaN	
6	2020-11-24	2	3	26.0	90.0	4232.0	
7	2020-11-25	3	3	NaN	96.0	NaN	
8	2020-11-26	2	2	27.0	100.0	2356.0	
9	2020-11-28	3	1	NaN	NaN	NaN	

10	2020-11-29	1	1	23.0	94.0	1254.0
11	2020-11-30	1	1	22.0	91.0	4232.0

	Product_Quality
0	A
1	A
2	A
3	NaN
4	B
5	B
6	C
7	NaN
8	B
9	A
10	A
11	A

```
[45]: sns.heatmap(df_saless.corr())
```

<ipython-input-45-3ad31eb899e5>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df_saless.corr())
```

```
[45]: <Axes: >
```



[45] :